

Inference speedup vs performance loss

