

Performance loss vs inference speedup

